
Self-Attentive Associative Memory

Hung Le¹ Truyen Tran¹ Svetha Venkatesh¹

Abstract

Heretofore, neural networks with external memory are restricted to single memory with lossy representations of memory interactions. A rich representation of relationships between memory pieces urges a high-order and segregated relational memory. In this paper, we propose to separate the storage of individual experiences (item memory) and their occurring relationships (relational memory). The idea is implemented through a novel Self-attentive Associative Memory (SAM) operator. Founded upon outer product, SAM forms a set of associative memories that represent the hypothetical high-order relationships between arbitrary pairs of memory elements, through which a relational memory is constructed from an item memory. The two memories are wired into a single sequential model capable of both memorization and relational reasoning. We achieve competitive results with our proposed two-memory model in a diversity of machine learning tasks, from challenging synthetic problems to practical testbeds such as geometry, graph, reinforcement learning, and question answering.

1. Introduction

Humans excel in remembering items and the relationship between them over time (Olson et al., 2006; Konkel & Cohen, 2009). Numerous neurocognitive studies have revealed this striking ability is largely attributed to the perirhinal cortex and hippocampus, two brain regions that support item memory (e.g., objects, events) and relational memory (e.g., locations of objects, orders of events), respectively (Cohen et al., 1997; Buckley, 2005). Relational memory theory posits that there exists a representation of critical relationships amongst arbitrary items, which allows inferential reasoning capacity (Eichenbaum, 1993; Zeithamova

et al., 2012). It remains unclear how the hippocampus can select the stored items in clever ways to unearth their hidden relationships and form the relational representation.

Research on artificial intelligence has focused on designing item-based memory models with recurrent neural networks (RNNs) (Hopfield, 1982; Elman, 1990; Hochreiter & Schmidhuber, 1997) and memory-augmented neural networks (MANNs) (Graves et al., 2014; 2016; Le et al., 2018a; 2019). These memories support long-term retrieval of previously seen items yet lack explicit mechanisms to represent arbitrary relationships amongst the constituent pieces of the memories. Recently, further attempts have been made to foster relational modeling by enabling memory-memory interactions, which is essential for relational reasoning tasks (Santoro et al., 2017; 2018; Vaswani et al., 2017). However, no effort has been made to model jointly item memory and relational memory explicitly.

We argue that dual memories in a single system are crucial for solving problems that require both memorization and relational reasoning. Consider graphs wherein each node is associated with versatile features— as example a road network structure where each node is associated with diverse features: *graph 1* where the nodes are building landmarks and *graph 2* where the nodes are flora details. The goal here is to reason over the structure and output the associated features of the nodes instead of the pointer or index to the nodes. Learning to output associated node features enables generalization to entirely novel features, i.e., a model can be trained to generate a navigation path with building landmarks (*graph 1*) and tested in the novel context of generating a navigation path with flora landmarks (*graph 2*). This may be achieved if the model stores the features and structures into its item and relational memory, separately, and reason over the two memories using rules acquired during training.

Another example requiring both item and relational memory can be understood by amalgamating the N^{th} -farthest (Santoro et al., 2018) and associative recall (Graves et al., 2014) tasks. N^{th} -farthest requires relational memory to return a fixed one-hot encoding representing the index to the N^{th} -farthest item, while associative recall returns the item itself, requiring item memory. If these tasks are amalgamated to compose Relational Associative Recall (RAR) – return the N^{th} -farthest item from a query (see § 3.2), it is clear that

¹Applied AI Institute, Deakin University, Geelong, Australia. Correspondence to: Hung Le <thai.le@deakin.edu.au>.

both item and relational memories are required.

Three limitations of the current approaches are: (i) the relational representation is often computed without storing, which prevents reusing the precomputed relationships in sequential tasks (Vaswani et al., 2017; Santoro et al., 2017), (ii) few works that manage both items and the relationships in a single memory, make it hard to understand how relational reasoning occurs (Santoro et al., 2018; Schlag & Schmidhuber, 2018), (iii) the memory-memory relationship is coarse since it is represented as either dot product attention (Vaswani et al., 2017) or weighted summation via neural networks (Santoro et al., 2017). Concretely, the former uses a scalar to measure cosine distance between two vectors and the latter packs all information into one vector via only additive interactions.

To overcome the current limitations, we hypothesize a two-memory model, in which the relational memory exists separately from the item memory. To maintain a rich representation of the relationship between items, the relational memory should be higher-order than the item memory. That is, the relational memory stores multiple relationships, each of which should be represented by a matrix rather than a scalar or vector. Otherwise, the capacity of the relational memory is downgraded to that of the item memory. Finally, as there are two separate memories, they must communicate to enrich the representation of one another.

To implement our hypotheses, we introduce a novel operator that facilitates the communication from the item memory to the relational memory. The operator, named Self-attentive Associative Memory (SAM) leverages the dot product attention with our outer product attention. Outer product is critical for constructing higher-order relational representations since it retains bit-level interactions between two input vectors, thus has potential for rich representational learning (Smolensky, 1990). SAM transforms a second-order (matrix) item memory into a third-order relational representation through two steps. First, SAM decodes a set of patterns from the item memory. Second, SAM associates each pair of patterns using outer product and sums them up to form a hetero-associative memory. The memory thus stores relationships between stored items accumulated across timesteps to form a relational memory.

The role of item memory is to memorize the input data over time. To selectively encode the input data, the item memory is implemented as a gated auto-associative memory. Together with previous read-out values from the relational memory, the item memory is used as the input for SAM to construct the relational memory. In return, the relational memory transfers its knowledge to the item memory through a distillation process. The backward transfer triggers recurrent dynamics between the two memories, which may be essential for simulating hippocampal processes (Kumaran

& McClelland, 2012). Another distillation process is used to transform the relational memory to the output value.

Taken together, we contribute a new neural memory model dubbed SAM-based Two-memory Model (STM) that takes inspiration from the existence of both item and relational memory in human brain (Konkel & Cohen, 2009). In this design, the relational memory is higher-order than the item memory and thus necessitates a core operator that manages the information exchange from the item memory to the relational memory. The operator, namely Self-attentive Associative Memory (SAM), utilizes outer product to construct a set of hetero-associative memories representing relationships between arbitrary stored items. We apply our model to a wide range of tasks that may require both item and relational memory: various algorithmic learning, geometric and graph reasoning, reinforcement learning and question-answering tasks. Several analytical studies on the characteristics of our proposed model are also given in the Appendix.

2. Methods

2.1. Outer product attention (OPA)

Outer product attention (OPA) is a natural extension of the query-key-value dot product attention (Vaswani et al., 2017). Dot product attention (DPA) for single query q and n_{kv} pairs of key-value can be formulated as follows,

$$A^\circ(q, K, V) = \sum_{i=1}^{n_{kv}} \mathcal{S}(q \cdot k_i) v_i \quad (1)$$

where $A^\circ \in \mathbb{R}^{d_v}$, $q, k_i \in \mathbb{R}^{d_{qk}}$, $v_i \in \mathbb{R}^{d_v}$, \cdot is dot product, and \mathcal{S} forms softmax function. We propose a new outer product attention with similar formulation yet different meaning,

$$A^\otimes(q, K, V) = \sum_{i=1}^{n_{kv}} \mathcal{F}(q \odot k_i) \otimes v_i \quad (2)$$

where $A^\otimes \in \mathbb{R}^{d_{qk} \times d_v}$, $q, k_i \in \mathbb{R}^{d_{qk}}$, $v \in \mathbb{R}^{d_v}$, \odot is element-wise multiplication, \otimes is outer product and \mathcal{F} is chosen as element-wise tanh function.

A crucial difference between DPA and OPA is that while the former retrieves an attended item A° , the latter forms a relational representation A^\otimes . As a relational representation, A^\otimes captures all bit-level associations between the key-scaled query and the value. This offers two benefits: (i) a higher-order representational capacity that DPA cannot provide and (ii) a form of associative memory that can be later used to retrieve stored item by using a contraction operation $\mathcal{P}(A^\otimes)$ (see Appendix § C-Prop. 6).

OPA is closely related to DPA. The relationship between the two for simple \mathcal{S} and \mathcal{F} is presented as follows,

Proposition 1. Assume that \mathcal{S} is a linear transformation: $\mathcal{S}(x) = ax + b$ ($a, b, x \in \mathbb{R}$), we can extract A° from A^\otimes by using an element-wise linear transformation $\mathcal{F}(x) = a^f \odot x + b^f$ ($a^f, b^f, x \in \mathbb{R}^{d_{qk}}$) and a contraction $\mathcal{P}: \mathbb{R}^{d_{qk} \times d_v} \rightarrow \mathbb{R}^{d_v}$ such that

$$A^\circ(q, K, V) = \mathcal{P}(A^\otimes(q, K, V)) \quad (3)$$

Proof. see Appendix § A. \square

Moreover, when $n_{kv} = 1$, applying a high dimensional transformation $\mathcal{G}(A^\otimes)$ is equivalent to the well-known bi-linear model (see Appendix § B-Prop. 4). By introducing OPA, we obtain a new building block that naturally supports both powerful relational bindings and item memorization.

2.2. Self-attentive Associative Memory (SAM)

We introduce a novel and generic operator based upon OPA that constructs relational representations from an item memory. The relational information is extracted via preserving the outer products between any pairs of items from the item memory. Hence, we name this operator Self-attentive Associative Memory (SAM). Given an item memory $M \in \mathbb{R}^{n \times d}$ and parametric weights $\theta = \{W_q \in \mathbb{R}^{n_q \times n}, W_k \in \mathbb{R}^{n_{kv} \times n}, W_v \in \mathbb{R}^{n_{kv} \times n}\}$, SAM retrieves n_q queries, n_{kv} keys and values from M as M_q, M_k and M_v , respectively,

$$M_q = \mathcal{LN}(W_q M) \quad (4)$$

$$M_k = \mathcal{LN}(W_k M) \quad (5)$$

$$M_v = \mathcal{LN}(W_v M) \quad (6)$$

where \mathcal{LN} is layer normalization operation (Ba et al., 2016b). Then SAM returns a relational representation $\text{SAM}_\theta(M) \in \mathbb{R}^{n_q \times d \times d}$, in which the s -th element of the first dimension is defined as

$$\text{SAM}_\theta(M)[s] = A^\otimes(M_q[s], M_k, M_v) \quad (7)$$

$$= \sum_{j=1}^{n_{kv}} \mathcal{F}(M_q[s] \odot M_k[j]) \otimes M_v[j] \quad (8)$$

where $s = 1, \dots, n_q$. $M_q[s]$, $M_k[j]$ and $M_v[j]$ denote the s -th row vector of matrix M_q , the j -th row vector of matrix M_k and M_v , respectively. A diagram illustrating SAM operations is given in Fig. 1 (right).

It should be noted that M can be any item memory including the slot-based memories (Le et al., 2019), direct inputs (Vaswani et al., 2017) or associative memories (Kohonen, 1972; Hopfield, 1982). We choose $M \in \mathbb{R}^{d \times d}$ as a form of classical associative memory, which is biologically plausible (Marr & Thach, 1991). Here, we follow the

traditional practice that sets $n = d$ for the associative item memory. From M we read query, key and value items to form $\text{SAM}_\theta(M)$ —a new set of hetero-associative memories using Eq. 8. Each hetero-associative memory represents the relationship between a query and all values. The role of the keys is to maintain possible perfect retrieval for the item memory (Appendix § C-Prop. 6).

The high-order structure of SAM allows it to preserve bit-level relationships between a query and a value in a matrix. SAM compresses several relationships with regard to a query by summing all the matrices to form a hetero-associative memory containing d^2 scalars, where d is the dimension of M . As there are n_{kv} relationships given 1 query, the summation results in on average d^2/n_{kv} scalars of representation per relationship, which is greater than 1 if $d > \sqrt{n_{kv}}$. By contrast, current self-attention mechanisms use dot product to measure the relationship between any pair of memory slots, which means 1 scalar per relationship.

2.3. SAM-based Two-Memory Model (STM)

To effectively utilize the SAM operator, we design a system which consists of two memory units $\mathcal{M}_t^i \in \mathbb{R}^{d \times d}$ and $\mathcal{M}_t^r \in \mathbb{R}^{n_q \times d \times d}$: one for items and the other for relationships, respectively. From a high-level view, at each timestep, we use the current input data x_t and the previous state of memories $\{\mathcal{M}_{t-1}^i, \mathcal{M}_{t-1}^r\}$ to produce output o_t and new state of memories $\{\mathcal{M}_t^i, \mathcal{M}_t^r\}$. The memory executions are described as follows.

\mathcal{M}_t^i -Write The item memory distributes the data from the input across its rows in the form of associative memory. For an input x_t , we update the item memory as

$$\begin{aligned} X_t &= f_1(x_t) \otimes f_2(x_t) \\ \mathcal{M}_t^i &= \mathcal{M}_{t-1}^i + X_t \end{aligned} \quad (9)$$

where f_1 and f_2 are feed-forward neural networks that output d -dimensional vectors. This update does not discriminate the input data and inherits the low-capacity of classical associative memory (Rojas, 2013). We leverage the gating mechanisms of LSTM (Hochreiter & Schmidhuber, 1997) to improve Eq. 9 as

$$\mathcal{M}_t^i = F_t(\mathcal{M}_{t-1}^i, x_t) \odot \mathcal{M}_{t-1}^i + I_t(\mathcal{M}_{t-1}^i, x_t) \odot X_t \quad (10)$$

where F_t and I_t are forget and input gates, respectively. Detailed implementation of these gates is in Appendix § D.

\mathcal{M}_t^r -Read As relationships stored in \mathcal{M}_t^r are represented as associative memories, the relational memory can be read to reconstruct previously seen items. As shown in Appendix § C-Prop. 7, the read is basically a two-step contraction,




Figure 1. STM (left) and SAM (right). SAM uses neural networks θ to extract query, key and value elements from a matrix memory M . In this illustration, $n_q = 3$ and $n_{kv} = 4$. Then, it applies outer product attention to output a 3D tensor relational representation. In STM, at every timestep, the item memory \mathcal{M}_t^i is updated with new input x_t using gating mechanisms (Eq. 10). The item memory plus the read-out from the relational memory is forwarded to SAM, resulting in a new relational representation to update the relational memory \mathcal{M}_t^r (Eq. 11-12). The relational memory transfers its knowledge to the item memory (Eq. 13) and output value (Eq. 14).

$$v_t^r = \text{softmax} \left(f_3(x_t)^\top \right) \mathcal{M}_{t-1}^r f_2(x_t) \quad (11)$$

where f_3 is a feed-forward neural network that outputs a n_q -dimensional vector. The read value provides an additional input coming from the previous state of \mathcal{M}^r to relational construction process, as shown later in Eq. 12.

\mathcal{M}^i -Read \mathcal{M}^r -Write We use SAM to read from \mathcal{M}^i and construct a candidate relational memory, which is simply added to the previous relational memory to perform the relational update,

$$\mathcal{M}_t^r = \mathcal{M}_{t-1}^r + \alpha_1 \text{SAM}_\theta \left(\mathcal{M}_t^i + \alpha_2 v_t^r \otimes f_2(x_t) \right) \quad (12)$$

where α_1 and α_2 are blending hyper-parameters. The input for SAM is a combination of the current item memory \mathcal{M}_t^i and the association between the extracted item from the previous relational memory v_t^r and the current input data x_t . Here, v_t^r enhances the relational memory with information from the distant past. The resulting relational memory stores associations between several pairs of items in a 3D tensors of size $n_q \times d \times d$. In our SAM implementation, $n_{kv} = n_q$.

\mathcal{M}^r -Transfer In this phase, the relational knowledge from \mathcal{M}_t^r is transferred to the item memory by using high dimensional transformation,

$$\mathcal{M}_t^i = \mathcal{M}_t^i + \alpha_3 \mathcal{G}_1 \circ \mathcal{V}_f \circ \mathcal{M}_t^r \quad (13)$$

where \mathcal{V}_f is a function that flattens the first two dimensions of its input tensor, \mathcal{G}_1 is a feed-forward neural network

that maps $\mathbb{R}^{(n_q d) \times d} \rightarrow \mathbb{R}^{d \times d}$ and α_3 is a blending hyper-parameter. As shown in Appendix § B-Prop. 5, with trivial \mathcal{G}_1 , the transfer behaves as if the item memory is enhanced with long-term stored values from the relational memory. Hence, \mathcal{M}^r -Transfer is also helpful in supporting long-term recall (empirical evidences in § 3.1). In addition, at each timestep, we distill the relational memory into an output vector $o_t \in \mathbb{R}^{n_o}$. We alternatively flatten and apply high-dimensional transformations as follow,

$$o_t = \mathcal{G}_3 \circ \mathcal{V}_l \circ \mathcal{G}_2 \circ \mathcal{V}_l \circ \mathcal{M}_t^r \quad (14)$$

where \mathcal{V}_l is a function that flattens the last two dimensions of its input tensor. \mathcal{G}_2 and \mathcal{G}_3 are two feed-forward neural networks that map $\mathbb{R}^{n_q \times (dd)} \rightarrow \mathbb{R}^{n_q \times n_r}$ and $\mathbb{R}^{n_q n_r} \rightarrow \mathbb{R}^{n_o}$, respectively. n_r is a hyper-parameter.

Unlike the contraction (Eq. 11), the distillation process does not simply reconstruct the stored items. Rather, thanks to high-dimensional transformations, it captures bi-linear representations stored in the relational memory (proof in Appendix § B). Hence, despite its vector form, the output of our model holds a rich representation that is useful for both sequential and relational learning. We discuss further on how to quantify the degree of relational distillation in Appendix § G. The summary of components of STM is presented in Fig. 1 (left).

3. Results

3.1. Ablation study

We test different model configurations on two classical tasks for sequential and relational learning: associative retrieval (Ba et al., 2016a) and N^{th} -farthest (Santoro et al., 2018)

Model	Length 30		Length 50	
	E.	A.	E.	A.
Fast weight*	50	100	5000	20.8
WeiNet*	35	100	50	100
STM ($d = 48$, w/o transfer)	10	100	100	100
STM ($d = 48$)	20	100	80	100
STM ($d = 96$, w/o gates)	100	24	100	20
STM ($d = 96$)	10	100	20	100

Table 1. Comparison of models on associative retrieval task with number of epochs E. required to converge (lower is better) and convergence test accuracy A. (% , higher is better). * is reported from Zhang & Zhou (2017).

(see Appendix § E for task details and learning curves). Our source code is available at <https://github.com/thaihungle/SAM>.

Associative retrieval This task measures the ability to recall a seen item given its associated key and thus involves item memory. We use the setting with input sequence length 30 and 50 (Zhang & Zhou, 2017). Three main factors affecting the item memory of STM are the dimension d of the auto-associative item memory, the gating mechanisms (Eq. 10) and the relational transfer (Eq. 13). Hence, we ablate our STM ($d = 96$, full features) by creating three other versions: small STM with transfer ($d = 48$), small STM without transfer ($d = 48$, w/o transfer) and STM without gates ($d = 96$, w/o gates). n_q is fixed to 1 as the task does not require much relational learning.

Table 1 reports the number of epochs required to converge and the final testing accuracy. Without the proposed gating mechanism, STM struggles to converge, which highlights the importance of extending the capacity of the auto-associative item memory. The convergence speed of STM is significantly improved with a bigger item memory size. Relational transfer seems more useful for longer input sequences since if requested, it can support long-term retrieval. Compared to other fast-weight baselines, the full-feature STM performs far better as it needs only 10 and 20 epochs to solve the tasks of length 30 and 50, respectively.

N^{th} -farthest This task evaluates the ability to learn the relationship between stored vectors. The goal is to find the N^{th} -farthest vector from a query vector, which requires a relational memory for distances between vectors and a sorting mechanism over the distances. For relational reasoning tasks, the pivot is the number of extracted items n_q for establishing the relational memory. Hence, we run our STM with different $n_q = 1, 4, 8$ using the same problem setting (8 16-dimensional input vectors), optimizer (Adam), batch size (1600) as in Santoro et al. (2018). We also run the task with TPR (Schlag & Schmidhuber, 2018)—a high-order

Model	Accuracy (%)
DNC*	25
RMC*	91
TPR	13
STM ($n_q = 1$)	84
STM ($n_q = 4$)	95
STM ($n_q = 8$)	98

Table 2. Comparison of models on N^{th} -farthest task (test accuracy). * is reported from Santoro et al. (2018).

fast-weight model that is designed for reasoning.

As reported in Table 2, increasing n_q gradually improves the accuracy of STM. As there are 8 input vectors in this task, literally, at each timestep the model needs to extract 8 items to compute all pairs of distances. However, as the extracted item is an entangled representation of all stored vectors and the temporarily computed distances are stored in separate high-order storage, even with $n_q = 1, 4$, STM achieves moderate results. With $n_q = 8$, STM nearly solves the task perfectly, outperforming RMC by a large margin. We have tried to tune TPR for this task without success (see Appendix § E). This illustrates the challenge of training high-order neural networks in diverse contexts.

3.2. Algorithmic synthetic tasks

Algorithmic synthetic tasks (Graves et al., 2014) examine sequential models on memorization capacity (eg., Copy, Associative recall) and simple relational reasoning (eg., Priority sort). Even without explicit relational memory, MANNs have demonstrated good performance (Graves et al., 2014; Le et al., 2020), but they are verified for only low-dimensional input vectors (< 8 bits). As higher-dimensional inputs necessitate higher-fidelity memory storage, we evaluate the high-fidelity reconstruction capacity of sequential models for these algorithmic tasks with 32-bit input vectors.

Two chosen algorithmic tasks are Copy and Priority sort. Item memory is enough for Copy where the models just output the input vectors seen in the same order in which they are presented. For Priority sort, a relational operation that compares the priority of input vectors is required to produce the seen input vectors in the sorted order according to the priority score attached to each input vector. The relationship is between input vectors and thus simply first-order (see Appendix § G for more on the order of relationship).

Inspired by Associative recall and N^{th} -farthest tasks, we create a new task named Relational Associative Recall (RAR). In RAR, the input sequence is a list of items followed by a query item. Each item is a list of several 32-bit vectors and thus can be interpreted as a concatenated long




Figure 2. Bit error per sequence vs training iteration for algorithmic synthetic tasks.

vector. The requirement is to reconstruct the seen item that is farthest or closest (yet unequal) to the query. The type of the relationship is conditioned on the last bit of the query vector, i.e., if the last bit is 1, the target is the farthest and 0 the closest. The evaluated models must compute the distances from the query item to any other seen items and then compare the distances to find the farthest/closest one. Hence, this task is similar to the N^{th} -farthest task, which is second-order relational and thus needs relational memory. However, this task is more challenging since the models must reconstruct the seen items (32-bit vectors). Compared to $N = 8$ possible one-hot outputs in N^{th} -farthest, the output space in RAR is 2^{32} per step, thereby requiring high-fidelity item memory.

We evaluate our model STM ($n_q = 8$, $d = 96$) with the 4 following baselines: LSTM (Hochreiter & Schmidhuber, 1997), attentional LSTM (Bahdanau et al., 2015), NTM (Graves et al., 2014) and RMC (Santoro et al., 2018). Details of the implementation are listed in Appendix § F. The learning curves (mean and error bar over 5 runs) are presented in Fig. 2.

LSTM is often the worst performer as it is based on vector memory. ALSTM is especially good for Copy as it has a privilege to access input vectors at every step of decoding. However, when dealing with relational reasoning, memory-less attention in ALSTM does not help much. NTM performs well on Copy and moderately on Priority sort, yet badly on RAR possibly due to its bias towards item memory. Although equipped with self-attention relational memory, RMC demonstrates trivial performance on all tasks. This suggests a limitation of using dot-product attention to represent relationships when the tasks stress memorization or the relational complexity goes beyond dot-product capacity. Amongst all models, only the proposed STM demonstrates consistently good performance where it almost achieves zero errors on these 3 tasks. Notably, for RAR, only STM can surpass the bottleneck error of 30 bits and reach ≈ 1 bit error, corresponding to 0% and 87% of items perfectly reconstructed, respectively.

3.3. Geometric and graph reasoning

Problems on geometry and graphs are a good testbed for relational reasoning, where geometry stipulates spatial relationships between points, and graphs the relational structure of nodes and edges. Classical problems include Convex hull, Traveling salesman problem (TSP) for geometry, and Shortest path, Minimum spanning tree for graph. Convex hull and TSP data are from Vinyals et al. (2015) where input sequence is a list of points’ coordinates (number of points $N \sim [5, 20]$). Graphs in Shortest path and Minimum spanning tree are generated with solutions found by Dijkstra and Kruskal algorithms, respectively. A graph input is represented as a sequence of triplets ($node_1, node_2, edge_{12}$). The desired output is a sequence of associated features of the solution points/nodes (more in Appendix § H).

We generate a random one-hot associated feature for each point/node, which is stacked into the input vector. This allows us to output the node’s associated features. This is unlike Vinyals et al. (2015), who just outputs the pointers to the nodes. Our modification creates a challenge for both training and testing. The training is more complex as the feature of the nodes varies even for the same graph. The testing is challenging as the associated features are likely to be different from that in the training. A correct prediction for a timestep is made when the predicted feature matches perfectly with the ground truth feature in the timestep. To measure the performance, we use the average accuracy of prediction across steps. We use the same baselines as in § 3.2 except that we replace NTM with DNC as DNC performs better on graph reasoning (Graves et al., 2016).

We report the best performance of the models on the testing datasets in Table 3. Although our STM has fewest parameters, it consistently outperforms other baselines by a significant margin. As usual, LSTM demonstrates an average performance across tasks. RMC and ALSTM are only good at Convex hull. DNC performs better on graph-like problems such as Shortest path and Minimum spanning tree. For the NP-hard TSP ($N = 5$), despite moderate point ac-

Model	#Parameters	Convex hull		TSP		Shortest path	Minimum spanning tree
		$N = 5$	$N = 10$	$N = 5$	$N = 10$		
LSTM	4.5 M	89.15	82.24	73.15 (2.06)	62.13 (3.19)	72.38	80.11
ALSTM	3.7 M	89.92	85.22	71.79 (2.05)	55.51 (3.21)	76.70	73.40
DNC	1.9 M	89.42	79.47	73.24 (2.05)	61.53 (3.17)	83.59	82.24
RMC	2.8 M	93.72	81.23	72.83 (2.05)	37.93 (3.79)	66.71	74.98
STM	1.9 M	96.85	91.88	73.96 (2.05)	69.43 (3.03)	93.43	94.77

Table 3. Prediction accuracy (%) for geometric and graph reasoning with random one-hot features. Italic numbers are tour length-additional metric for TSP. Average optimal tour lengths found by brute-force search for $N = 5$ and 10 are 2.05 and 2.88, respectively.




Figure 3. Average reward vs number of games for reinforcement learning task in n -frame skip settings.

curacy, all models achieve nearly minimal solutions with an average tour length of 2.05. When increasing the difficulty with more points ($N = 10$), none of these models reach an average optimal tour length of 2.88. However, only STM approaches closer to the optimal solution without the need for pointer and beam search mechanisms. Armed with both item and relational memory, STM’s superior performance suggests a qualitative difference in the way STM and other methods solve these problems.

3.4. Reinforcement learning

Memory is helpful for partially observable Markov decision process (Bakker, 2002). We apply our memory to LSTM agents in Atari game environment using A3C training (Mnih et al., 2016). More details are given in Appendix § I. In Atari games, each state is represented as the visual features of a video frame and thus is partially observable. To perform well, RL agents should remember and relate several frames to model the game state comprehensively. These abilities are challenged when over-sampling and under-sampling the observation, respectively. We analyze the performance of LSTM agents and their STM-augmented counterparts under these settings using a game: Pong.

To be specific, we test the two agents on different frame skips (0, 4, 16, 32). We create n -frame skip setting by allowing the agent to see the environment only after every n frames, where 4-frame skip is standard in most Atari environments. When no frameskip is applied (over-sampling), the number of observations is dense and the game is long (up to 9000 steps per game), which requires high-capacity

item memory. On the contrary, when a lot of frames are skipped (under-sampling), the observations become scarce and the agents must model the connection between frames meticulously, demanding better relational memory.

We run each configuration 5 times and report the mean and error bar of moving average reward (window size = 100) through training time in Fig. 3. In a standard condition (4-frame skip), both baselines can achieve perfect performance and STM outperforms LSTM slightly in terms of convergence speed. The performance gain becomes clearer under extreme conditions with over-sampling and under-sampling. STM agents require fewer practices to accomplish higher rewards, especially in the 32-frame skip environment, which illustrates that having strong item and relational memory in a single model is beneficial to RL agents.

3.5. Question answering

bAbI is a question answering dataset that evaluates the ability to remember and reason on textual information (Weston et al., 2015). Although synthetically generated, the dataset contains 20 challenging tasks such as pathfinding and basic induction, which possibly require both item and relational memory. Following Schlag & Schmidhuber (2018), each story is preprocessed into a sentence-level sequence, which is fed into our STM as the input sequence. We jointly train STM for all tasks using normal supervised training (more in Appendix § J). We compare our model with recent memory networks and report the results in Table 4.

MANNs such as DNC and NUTM have strong item memory, yet do not explicitly support relational learning, leading to

Model	Error	
	Mean	Best
DNC (Graves et al., 2016)	12.8 ± 4.7	3.8
NUTM (Le et al., 2020)	5.6 ± 1.9	3.3
TPR (Schlag & Schmidhuber, 2018)	1.34 ± 0.52	0.81
UT (Dehghani et al., 2018)	1.12 ± 1.62	0.21
MNM-p (Munkhdalai et al., 2019)	0.55 ± 0.74	0.18
STM	0.39 ± 0.18	0.15

Table 4. bAbI task: mean \pm std. and best error over 10 runs.

significantly higher errors compared to other models. On the contrary, TPR is explicitly equipped with relational bindings but lack of item memory and thus clearly underperforms our STM. Universal Transformer (UT) supports a manually set item memory with dot product attention, showing higher mean error than STM with learned item memory and outer product attention. Moreover, our STM using normal supervised loss outperforms MNM-p trained with meta-level loss, establishing new state-of-the-arts on bAbI dataset. Notably, STM achieves this result with low variance, solving 20 tasks for 9/10 run (see Appendix § J).

4. Related Work

Background on associative memory Associative memory is a classical concept to model memory in the brain (Marr & Thach, 1991). While outer product is one common way to form the associative memory, different models employ different memory retrieval mechanisms. For example, Correlation Matrix Memory (CMM) and Hopfield network use dot product and recurrent networks, respectively (Kohonen, 1972; Hopfield, 1982). The distinction between our model and other associative memories lies in the fact that our model’s association comes from several pieces of the memory itself rather than the input data. Also, unlike other two-memory systems (Le et al., 2018b; 2020) that simulate data/program memory in computer architecture, our STM resembles item and relational memory in human cognition.

Background on attention Attention is a mechanism that allows interactions between a query and a set of stored keys/values (Graves et al., 2014; Bahdanau et al., 2014). Self-attention mechanism allows stored items to interact with each other either in forms of feed-forward (Vaswani et al., 2017) or recurrent (Santoro et al., 2018; Le et al., 2019) networks. Modeling memory interactions can also be achieved via attention over a set of parallel RNNs (Henaff et al., 2016). Although some form of relational memory can be kept in these approaches, they all use dot product attention to measure interactions per attention head as a scalar, and thus lose much relational information. We use outer product to represent the interactions as a matrix and thus our outer product self-attention is supposed to be richer

than the current self-attention mechanisms (Prop. 1).

SAM as fast-weight Outer product represents Hebbian learning—a fast learning rule that can be used to build fast-weights (von der Malsburg, 1981). As the name implies, fast-weights update whenever an input is introduced to the network and stores the input pattern temporarily for sequential processing (Ba et al., 2016a). Meta-trained fast-weights (Munkhdalai et al., 2019) and gating of fast-weights (Schlag & Schmidhuber, 2017; Zhang & Zhou, 2017) are introduced to improve memory capacity. Unlike these fast-weight approaches, our model is not built on top of other RNNs. Recurrency is naturally supported within STM.

The tensor product representation (TPR), which is a form of high-order fast-weight, can be designed for structural reasoning (Smolensky, 1990). In a recent work (Schlag & Schmidhuber, 2018), a third-order TPR resembles our relational memory \mathcal{M}_t^r where both are 3D tensors. However, TPR does not enable interactions amongst stored patterns through self-attention mechanism. The meaning of each dimension of the TPR is not related to that of \mathcal{M}_t^r . More importantly, TPR is restricted to question answering task.

SAM as bi-linear model Bi-linear pooling produces output from two input vectors by considering all pairwise bit interactions and thus can be implemented by means of outer product (Tenenbaum & Freeman, 2000). To reduce computation cost, either low-rank factorization (Yu et al., 2017) or outer product approximation (Pham & Pagh, 2013) is used. These approaches aim to enrich feed-forward layers with bi-linear poolings yet have not focused on maintaining a rich memory of relationships.

Low-rank bi-linear pooling is extended to perform visual attentions (Kim et al., 2018). It results in different formulation from our outer product attention, which is equivalent to full rank bi-linear pooling (§ 2.1). These methods are designed for static visual question answering while our approach is used to maintain a relational memory over time, which can be applied to any sequential problem.

5. Conclusions

We have introduced the SAM-based Two-memory Model (STM) that implements both item and relational memory. To wire up the two memory system, we employ a novel operator named Self-attentive Associative Memory (SAM) that constructs the relational memory from outer-product relationships between arbitrary pieces of the item memory. We apply read, write and transfer operators to access, update and distill the knowledge from the two memories. The ability to remember items and their relationships of the proposed STM is validated through a suite of diverse tasks including associative retrieval, N^{th} -farthest, vector

algorithms, geometric and graph reasoning, reinforcement learning and question answering. In all scenarios, our model demonstrates strong performance, confirming the usefulness of having both item and relational memory in one model.

ACKNOWLEDGMENTS

This research was partially funded by the Australian Government through the Australian Research Council (ARC). Prof Venkatesh is the recipient of an ARC Australian Laureate Fellowship (FL170100006).

References

- Ba, J., Hinton, G. E., Mnih, V., Leibo, J. Z., and Ionescu, C. Using fast weights to attend to the recent past. In *Advances in Neural Information Processing Systems*, pp. 4331–4339, 2016a.
- Ba, J. L., Kiros, J. R., and Hinton, G. E. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016b.
- Bahdanau, D., Cho, K., and Bengio, Y. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014.
- Bahdanau, D., Cho, K., and Bengio, Y. Neural machine translation by jointly learning to align and translate. *Proceedings of the International Conference on Learning Representations*, 2015.
- Bakker, B. Reinforcement learning with long short-term memory. In *Advances in neural information processing systems*, pp. 1475–1482, 2002.
- Buckley, M. J. The role of the perirhinal cortex and hippocampus in learning, memory, and perception. *The Quarterly Journal of Experimental Psychology Section B*, 58(3-4):246–268, 2005.
- Cohen, N. J., Poldrack, R. A., and Eichenbaum, H. Memory for items and memory for relations in the procedural/declarative memory framework. *Memory*, 5(1-2):131–178, 1997.
- Dehghani, M., Gouws, S., Vinyals, O., Uszkoreit, J., and Kaiser, Ł. Universal transformers. *arXiv preprint arXiv:1807.03819*, 2018.
- Eichenbaum, H. *Memory, amnesia, and the hippocampal system*. MIT press, 1993.
- Elman, J. L. Finding structure in time. *Cognitive science*, 14(2):179–211, 1990.
- Graves, A., Wayne, G., and Danihelka, I. Neural turing machines. *arXiv preprint arXiv:1410.5401*, 2014.
- Graves, A., Wayne, G., Reynolds, M., Harley, T., Danihelka, I., Grabska-Barwińska, A., Colmenarejo, S. G., Grefenstette, E., Ramalho, T., Agapiou, J., et al. Hybrid computing using a neural network with dynamic external memory. *Nature*, 538(7626):471–476, 2016.
- Henaff, M., Weston, J., Szlam, A., Bordes, A., and LeCun, Y. Tracking the world state with recurrent entity networks. *arXiv preprint arXiv:1612.03969*, 2016.
- Hochreiter, S. and Schmidhuber, J. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- Hopfield, J. J. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the national academy of sciences*, 79(8):2554–2558, 1982.
- Kim, J.-H., Jun, J., and Zhang, B.-T. Bilinear attention networks. In *Advances in Neural Information Processing Systems*, pp. 1564–1574, 2018.
- Kohonen, T. Correlation matrix memories. *IEEE transactions on computers*, 100(4):353–359, 1972.
- Konkel, A. and Cohen, N. J. Relational memory and the hippocampus: representations and methods. *Frontiers in neuroscience*, 3:23, 2009.
- Kumaran, D. and McClelland, J. L. Generalization through the recurrent interaction of episodic memories: a model of the hippocampal system. *Psychological review*, 119(3):573, 2012.
- Le, H., Tran, T., Nguyen, T., and Venkatesh, S. Variational memory encoder-decoder. In *Advances in Neural Information Processing Systems*, pp. 1508–1518, 2018a.
- Le, H., Tran, T., and Venkatesh, S. Dual memory neural computer for asynchronous two-view sequential learning. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery; Data Mining, KDD '18*, pp. 1637–1645, New York, NY, USA, 2018b. ACM. ISBN 978-1-4503-5552-0. doi: 10.1145/3219819.3219981. URL <http://doi.acm.org/10.1145/3219819.3219981>.
- Le, H., Tran, T., and Venkatesh, S. Learning to remember more with less memorization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=r1xlvi0qYm>.
- Le, H., Tran, T., and Venkatesh, S. Neural stored-program memory. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=rkxxA24FDr>.

- Marr, D. and Thach, W. T. A theory of cerebellar cortex. In *From the Retina to the Neocortex*, pp. 11–50. Springer, 1991.
- Mnih, V., Badia, A. P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Silver, D., and Kavukcuoglu, K. Asynchronous methods for deep reinforcement learning. In *International conference on machine learning*, pp. 1928–1937, 2016.
- Munkhdalai, T., Sordoni, A., Wang, T., and Trischler, A. Metalearned neural memory. In *Advances in Neural Information Processing Systems*, pp. 13310–13321, 2019.
- Olson, I. R., Page, K., Moore, K. S., Chatterjee, A., and Verfaellie, M. Working memory for conjunctions relies on the medial temporal lobe. *Journal of Neuroscience*, 26(17):4596–4601, 2006.
- Pham, N. and Pagh, R. Fast and scalable polynomial kernels via explicit feature maps. In *Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 239–247. ACM, 2013.
- Rojas, R. *Neural networks: a systematic introduction*. Springer Science & Business Media, 2013.
- Rudelson, M. and Vershynin, R. Sampling from large matrices: An approach through geometric functional analysis. *Journal of the ACM (JACM)*, 54(4):21, 2007.
- Santoro, A., Raposo, D., Barrett, D. G., Malinowski, M., Pascanu, R., Battaglia, P., and Lillicrap, T. A simple neural network module for relational reasoning. In *Advances in neural information processing systems*, pp. 4967–4976, 2017.
- Santoro, A., Faulkner, R., Raposo, D., Rae, J., Chrzanowski, M., Weber, T., Wierstra, D., Vinyals, O., Pascanu, R., and Lillicrap, T. Relational recurrent neural networks. In *Advances in Neural Information Processing Systems*, pp. 7299–7310, 2018.
- Schlag, I. and Schmidhuber, J. Gated fast weights for on-the-fly neural program generation. In *NIPS Metalearning Workshop*, 2017.
- Schlag, I. and Schmidhuber, J. Learning to reason with third order tensor products. In *Advances in Neural Information Processing Systems*, pp. 9981–9993, 2018.
- Smolensky, P. Tensor product variable binding and the representation of symbolic structures in connectionist systems. *Artificial intelligence*, 46(1-2):159–216, 1990.
- Tenenbaum, J. B. and Freeman, W. T. Separating style and content with bilinear models. *Neural computation*, 12(6):1247–1283, 2000.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. In *Advances in neural information processing systems*, pp. 5998–6008, 2017.
- Vinyals, O., Fortunato, M., and Jaitly, N. Pointer networks. In *Advances in Neural Information Processing Systems*, pp. 2692–2700, 2015.
- von der Malsburg, C. The correlation theory of brain function, 1981. URL <http://cogprints.org/1380/>.
- Weston, J., Bordes, A., Chopra, S., Rush, A. M., van Merriënboer, B., Joulin, A., and Mikolov, T. Towards ai-complete question answering: A set of prerequisite toy tasks. *arXiv preprint arXiv:1502.05698*, 2015.
- Yu, Z., Yu, J., Fan, J., and Tao, D. Multi-modal factorized bilinear pooling with co-attention learning for visual question answering. In *Proceedings of the IEEE international conference on computer vision*, pp. 1821–1830, 2017.
- Zeithamova, D., Schlichting, M. L., and Preston, A. R. The hippocampus and inferential reasoning: building memories to navigate future decisions. *Frontiers in human neuroscience*, 6:70, 2012.
- Zhang, W. and Zhou, B. Learning to update auto-associative memory in recurrent neural networks for improving sequence memorization. *ArXiv*, abs/1709.06493, 2017.